

DATA606 Project Update #1

AI-enhanced Credit Card Transaction Data Analysis

Omkar Pardeshi, Ajay Singh, Ethan Pollock

Where We Left From

- Dataset including 1.3 million rows of credit card transaction data
- Dates range from December 31, 2018 to June 21, 2020
- Variables include:
 - Credit card number
 - Merchant/store
 - Transaction category
 - Transaction amount
 - First name, Last name
 - Gender of card holder
 - Address details of cardholder
- Models selected to test
 - Random Forest
 - Support Vector Machines
 - Deep Learning Neural Networks
 - Gradient Boosted Trees
 - Specifically XGBoost
 - Clustering and Unsupervised Learning
 - Specifically K-Mean and Hierarchical Clustering

Feature	Traditional Fraud Detection	AI & ML Fraud Detection
Core Approach	Rule-based, relies on human expertise	Automated, driven by machine learning models and pattern recognition
Advantages	<ul style="list-style-type: none"> <input checked="" type="checkbox"/> Human oversight <input checked="" type="checkbox"/> Simple to implement 	<ul style="list-style-type: none"> <input checked="" type="checkbox"/> Real-time analysis <input checked="" type="checkbox"/> Adaptive learning <input checked="" type="checkbox"/> Detects non-linear patterns
Disadvantages	<ul style="list-style-type: none"> <input checked="" type="checkbox"/> High false positives <input checked="" type="checkbox"/> Limited scalability <input checked="" type="checkbox"/> Only detects linear patterns 	<ul style="list-style-type: none"> <input checked="" type="checkbox"/> Data dependent <input checked="" type="checkbox"/> Complex to implement
Response Time	Slow, after manual review	Instant, real-time scoring with generative AI tools
Scalability	Limited, needs constant manual rule updates	Scalable and self-improving over time
Accuracy & Adaptability	Low – static rules can't evolve quickly	Highly adaptable to new fraud detection solutions
Use in Business	Traditional banking or small-scale platforms	Used in fintech , e-wallets , and modern digital payment services

Progress To Date

- ❑ Package Installation and Imports
- ❑ Data Cleaning
- ❑ Feature Engineering
- ❑ Clustering
- ❑ Predictive Modeling
- ❑ Explainability
- ❑ Interactive EDA Dashboard
- ❑ PDF Report Generation

Package Installation and Imports

- *pip* automatically installs any missing packages
 - scikit-learn
 - pandas
 - matplotlib
 - seaborn
 - shap
 - joblib
 - xgboost

```
# ## 0) Install required packages (if missing)
import sys, subprocess, pkg_resources
required = ['scikit-learn','pandas','matplotlib','seaborn','shap','joblib','xgboost']
installed = {pkg.key for pkg in pkg_resources.working_set}
to_install = [p for p in required if p not in installed]
if to_install:
    print("Installing:", to_install)
    subprocess.check_call([sys.executable, '-m', 'pip', 'install'] + to_install)
else:
    print("All required packages already installed.")
```

```
All required packages already installed.
/tmp/ipython-input-2972911722.py:2: DeprecationWarning: pkg_resources is deprecated as
import sys, subprocess, pkg_resources
```

Package Installation and Imports

- Data Handling & Visualization
 - pandas
 - numpy
 - matplotlib
- Modeling & Evaluation
 - RandomForestClassifier
 - K-Means
 - classification_report
- Preprocessing & Pipelines
 - StandardScalar
 - OneHotEncoder
 - Pipeline
 - ColumnTransformer
- Utility & Explainability
 - joblib
 - shap

```
# ## 1) Standard imports
import pandas as pd, numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier, IsolationForest
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.cluster import KMeans
from sklearn.metrics import classification_report, roc_auc_score, confusion_matrix
import joblib
import shap
print('OK')
```

OK

Data Cleaning

- Displays:
 - Data types
 - Non-null counts
 - Memory usage
- Check for missing values across all columns
- Provides summary statistics for the *amt* column
- Show class distribution for the *is_fraud* target variable

```
# EDA
print(df.info())
print('\nMissing values per column:\n', df.isnull().sum())
if 'amt' in df.columns:
    display(df['amt'].describe())
if 'is_fraud' in df.columns:
    display(df['is_fraud'].value_counts(dropna=False))

10 city           1296675 non-null   object
11 state          1296675 non-null   object
12 zip            1296675 non-null   int64
13 lat             1296675 non-null   float64
14 long            1296675 non-null   float64
15 city_pop        1296675 non-null   int64
16 job             1296675 non-null   object
17 dob             1296675 non-null   object
18 trans_num       1296675 non-null   object
19 unix_time       1296675 non-null   int64
20 merch_lat       1296675 non-null   float64
21 merch_long      1296675 non-null   float64
22 is_fraud        1296675 non-null   int64
23 merch_zipcode   1100702 non-null   float64
dtypes: float64(6), int64(6), object(12)
memory usage: 237.4+ MB
None
```

Missing values per column:			
Unnamed: 0	0	trans_date_trans_time	0
cc_num	0	merchant	0
category	0	amt	0
first	0	last	0
gender	0	street	0
city	0	state	0
zip	0	lat	0
long	0	city_pop	0
job	0	dob	0
trans_num	0	unix_time	0
merch_lat	0	merch_long	0
merch_long	0	is_fraud	0
merch_zipcode	195973		
dtype: int64			
count	1.296675e+06		
mean	7.035104e+01		
std	1.603160e+02		
min	1.000000e+00		
25%	9.650000e+00		
50%	4.752000e+01		
75%	8.314000e+01		
max	2.894890e+04		
Name: amt, dtype: float64			
is_fraud			
0	1289169		
1	7506		
Name: count, dtype: int64			

Feature Engineering

- Execute key transformations
 - Datetime Parsing
 - Extract *hour*, *day*, and *month* from *trans_date_trans_time*
 - Monetary Feature
 - Applies log transformation from *amt* to *amt_log1p* to reduce skew
 - Age Calculation
 - Derives age in years from *dob*
- Adds temporal, scaled, and demographic features that can capture behavioral pattern
 - Important in fraud detection

```
# Feature engineering
df2 = df.copy()

# Parse datetime if present
if 'trans_date_trans_time' in df2.columns:
    df2['trans_date_trans_time'] = pd.to_datetime(df2['trans_date_trans_time'], errors='coerce')
    df2['trans_hour'] = df2['trans_date_trans_time'].dt.hour
    df2['trans_day'] = df2['trans_date_trans_time'].dt.day
    df2['trans_month'] = df2['trans_date_trans_time'].dt.month

# Amount features
if 'amt' in df2.columns:
    df2['amt_log1p'] = np.log1p(df2['amt']).astype(float)

# Age from dob if exists
if 'dob' in df2.columns:
    df2['dob'] = pd.to_datetime(df2['dob'], errors='coerce')
    df2['age'] = ((pd.Timestamp.now() - df2['dob']).dt.days / 365.25).fillna(0).astype(int)

display(df2.head())
```

Unnamed:	trans_date_trans_time	cc_num	merchant	category	amt	first	last	gender	street	...	unix_time	merch_lat	merch_long	is_fraud	merch_zipcode	trans_hour	trans_day	trans_month	amt_log1p	
0	0	2019-01-01 00:00:18	2703186189650295	fraud_Ripon Kub and Mann	misc_net	4.97	Jennifer	Banks	F	561	...	1325376018	36.011293	-82.048315	0	28705.0	0	1	1	1.78674
1	1	2019-01-01 00:00:44	63042337322	fraud_Heller Gutmann and Zieme	grocery_pos	107.23	Stephanie	Gill	F	561	...	1325376044	49.159047	-118.186462	0	NaN	0	1	1	4.68425
2	2	2019-01-01 00:00:51	38859492057661	fraud_Lind-Buckridge	entertainment	220.11	Edward	Sanchez	M	561	...	1325376051	43.150704	-112.154481	0	83236.0	0	1	1	5.39866

Clustering

- Group similar transaction and/or users using K-Means clustering
- Scales selected numeric features with StandardScalar
- Applies K-Means clustering
 - 2-4 clusters based on data size
- Assigns each record a cluster label
- Unsupervised segmentation can help uncover hidden patterns
 - E.g., high-risk transaction clusters

```
# Select numeric features for clustering (fallback if columns missing)
clust_features = [c for c in ['amt', 'amt_log1p', 'trans_hour', 'city_pop', 'age'] if c in df2.columns]
print('Clustering on:', clust_features)

if clust_features:
    Xc = df2[clust_features].fillna(0).astype(float)
    scaler = StandardScaler()
    Xs = scaler.fit_transform(Xc)
    kmeans = KMeans(n_clusters=min(4, max(2, len(df2)//5)), random_state=42)
    df2['cluster'] = kmeans.fit_predict(Xs)
    print('Cluster counts:')
    display(df2['cluster'].value_counts())
else:
    print('No clustering features found.')
```

```
Clustering on: ['amt', 'amt_log1p', 'trans_hour', 'city_pop', 'age']
Cluster counts:
cluster
0    431275
3    418567
1    404949
2    41884
Name: count, dtype: int64
```

Anomaly Detection

- Detect potentially fraudulent transactions using unsupervised learning
- Use IsolationForest that allows 1% contamination threshold
- Output
 - Adds *anomaly_score* as -1 for anomaly and 1 for normal
 - Adds binary *anomaly* flag for easy filtering
- Detects rare and suspicious behavior without requiring labeled fraud data

```
# Anomaly detection with IsolationForest
ad_features = [c for c in ['amt','amt_log1p','trans_hour','city_pop'] if c in df2.columns]
if ad_features:
    Xad = df2[ad_features].fillna(0).astype(float)
    iso = IsolationForest(n_estimators=100, contamination=0.01, random_state=42)
    df2['anomaly_score'] = iso.fit_predict(Xad) # -1 anomaly, 1 normal
    df2['anomaly'] = df2['anomaly_score'] == -1
    display(df2[['amt'] + [c for c in ad_features if c != 'amt'] + ['anomaly']].head(8))
else:
    print('No features available for anomaly detection.')
```

	amt	amt_log1p	trans_hour	city_pop	anomaly
0	4.97	1.786747	0	3495	False
1	107.23	4.684259	0	149	False
2	220.11	5.398660	0	4154	False
3	45.00	3.828641	0	1939	False
4	41.96	3.760269	0	99	False
5	94.63	4.560487	0	2158	False
6	44.54	3.818591	0	2691	False
7	71.65	4.285653	0	6018	False

Predictive Modeling

- Train a supervised learning classifier using labeled data (*is_fraud*)
- Pipeline components
 - Preprocessing: scales numerics and applies one-hot encoding to categoricals
 - Numerical: *amt*, *trans_hour*, *city_pop*, *age*
 - Categorical: *category*, *merchant*, *job*, *state*
 - Classifier:
RandomForestClassifier with 100 trees

```
# Prepare data for supervised modeling
if 'is_fraud' in df2.columns:
    target = 'is_fraud'
# Features selection: numeric + some categorical (limited)
numeric_feats = [c for c in ['amt','amt_log1p','trans_hour','trans_day','trans_month','city_pop','age'] if c in df2.columns]
cat_feats = [c for c in ['category','merchant','job','state'] if c in df2.columns]
X = df2[numeric_feats + cat_feats].copy()
y = df2[target].astype(int).fillna(0)

# Simple preprocessing pipeline
numeric_transformer = Pipeline(steps=[('scaler', StandardScaler())])
from sklearn import __version__ as sklver
ohe_kwargs = {'handle_unknown': 'ignore'}
# use correct arg name depending on sklearn version
if int(sklver.split('.')[1]) >= 4:
    ohe_kwargs['sparse_output'] = False
else:
    ohe_kwargs['sparse'] = False

cat_transformer = Pipeline(steps=[('ohe', OneHotEncoder(**ohe_kwargs))])

preprocessor = ColumnTransformer(transformers=[
    ('num', numeric_transformer, numeric_feats),
    ('cat', cat_transformer, cat_feats)
], remainder='drop')

clf = Pipeline(steps=[('pre', preprocessor),
                     ('clf', RandomForestClassifier(n_estimators=100, random_state=42))])
```

Predictive Modeling

- Data is split into train/test sets (stratified)
- Trains and evaluates pipelines:
 - classification_report
 - ROC AUC
- Saves final model to *fraud_model.joblib*
- Modular and version-safe pipeline ensures reliable, repeatable training with performance reporting

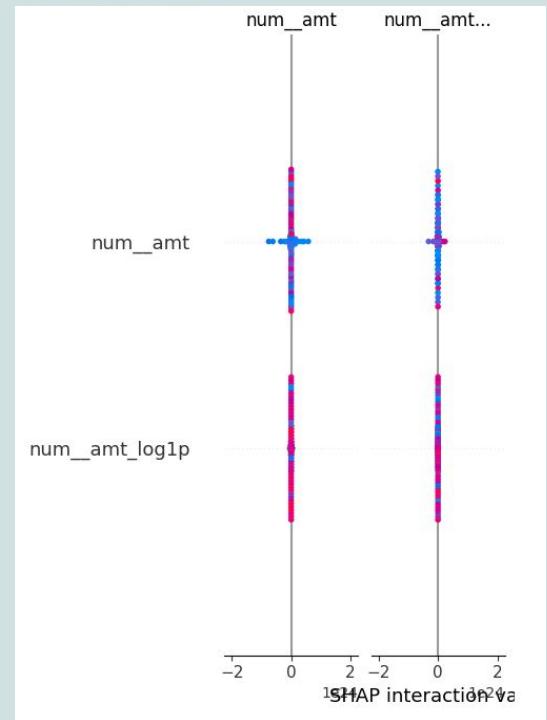
```
# train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y if len(y.unique())>1 else None)
clf.fit(X_train, y_train)
ypred = clf.predict(X_test)
print('Classification report:')
print(classification_report(y_test, ypred, zero_division=0))
if hasattr(clf[-1], 'predict_proba') and len(y_test)>0:
    yproba = clf.predict_proba(X_test)[:,1]
    try:
        print('ROC AUC:', roc_auc_score(y_test, yproba))
    except Exception as e:
        print('ROC AUC could not be computed:', e)
# Save model
joblib.dump(clf, 'fraud_model.joblib')
print('Model saved to fraud_model.joblib')
else:
    print('Target column `is_fraud` not found - skipping supervised modeling.')

Classification report:
          precision    recall  f1-score   support
           0       1.00     1.00    1.00    257934
           1       0.97     0.71    0.82     1501
   accuracy                           1.00    259335
  macro avg       0.99     0.86    0.91    259335
weighted avg       1.00     1.00    1.00    259335

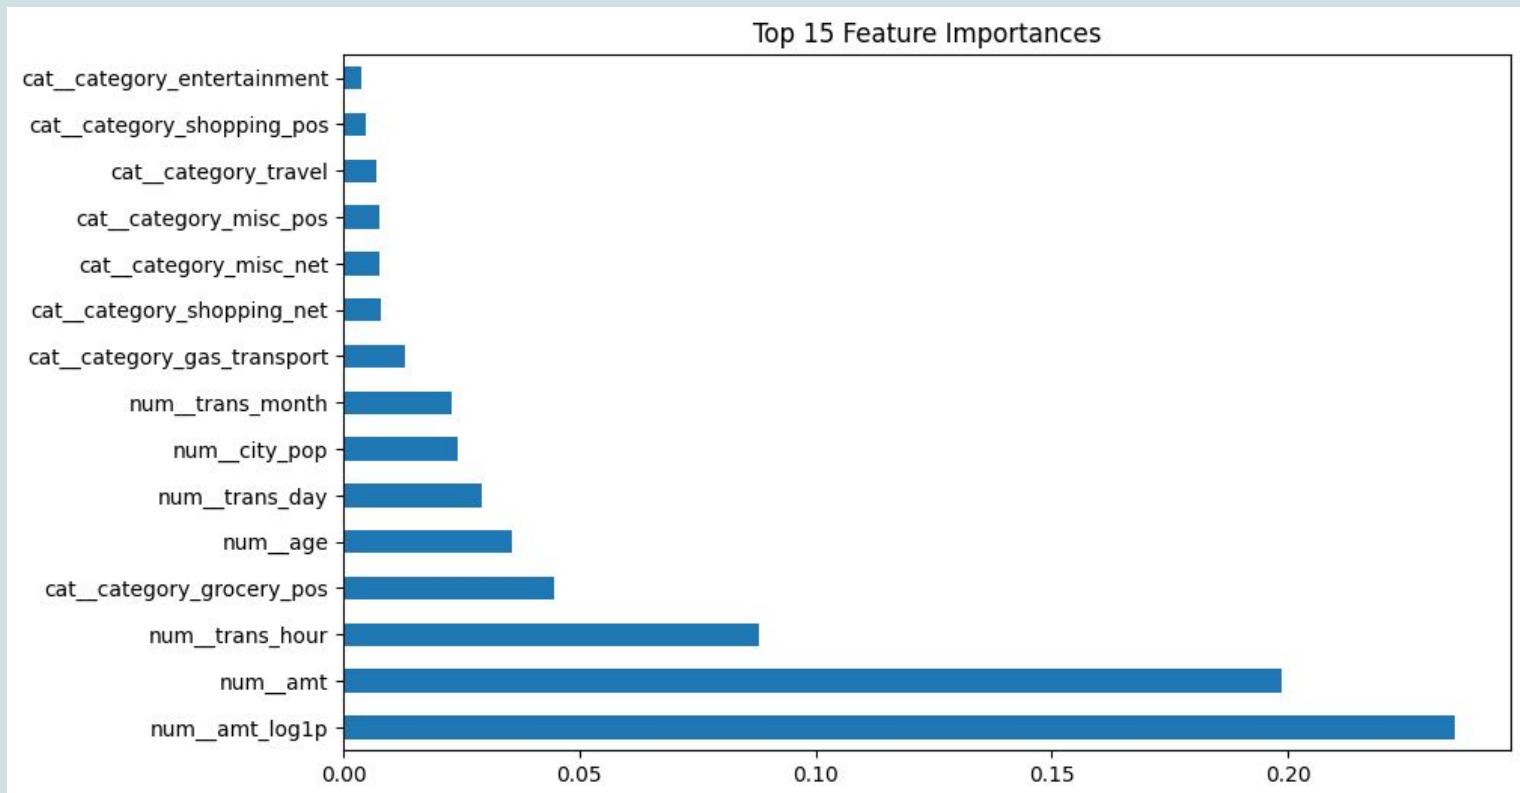
ROC AUC: 0.9911300629380465
Model saved to fraud_model.joblib
```

Explainability

- Automatically adjusts for SHAP API changes (legacy to modern)
- SHAP Summary Plot:
 - Uses top 100 transformed training samples for clarity
 - Visualizes feature impact and direction on model predictions
- Adds transparency to model decisions, build trust, and aids debugging or stakeholder communication



Explainability



Explainability

- Models tested:
 - RandomForestClassifier
 - GradientBoostingClassifier
 - LogisticRegression
- Reuses preprocessing pipeline for fair comparison
- Measures performance using ROC AUC on test data
- Outputs sorted table of model names against ROC AUC scores
- Helps identify the best-performing baseline model quickly without manual training

```
# === Cell 16: AutoML-style quick model comparison ===
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import roc_auc_score

models = {
    "RandomForest": RandomForestClassifier(n_estimators=100, random_state=42),
    "GradientBoosting": GradientBoostingClassifier(random_state=42),
    "LogisticRegression": LogisticRegression(max_iter=1000)
}

results = []
for name, m in models.items():
    pipe = Pipeline(steps=[('pre', preproc), ('clf', m)])
    pipe.fit(X_train, y_train)
    ypred = pipe.predict(X_test)
    yprob = pipe.predict_proba(X_test)[:,1] if hasattr(pipe, 'predict_proba') else ypred
    auc = roc_auc_score(y_test, yprob) if len(np.unique(y_test)) > 1 else np.nan
    results.append({'Model': name, 'ROC_AUC': auc})

pd.DataFrame(results).sort_values('ROC_AUC', ascending=False)
```

	Model	ROC_AUC
0	RandomForest	0.991130
1	GradientBoosting	0.971749
2	LogisticRegression	0.848463

Interactive EDA Dashboard

- Intuitive and interactive visualizations to explore fraud patterns in dataset
 - Histogram: distribution of transaction amounts, colored by fraud flag
 - Choropleth Map: fraud frequency by US State (geographic heatmap)
 - Bar Chart: average fraud rate by transaction category
- Plotly.express allows for dynamic and interactive plotting
- Enables stakeholders to explore trends, hotspots, and anomalies visually
 - Ideal for presentations and dashboards to stakeholders

```
# === Cell 17: Interactive EDA Dashboard (Plotly) ===
import plotly.express as px

if 'amt' in df2.columns and 'category' in df2.columns:
    fig1 = px.histogram(df2, x='amt', color='is_fraud', nbins=40,
                        title='Transaction Amount Distribution by Fraud Flag')
    fig1.show()

if 'state' in df2.columns:
    fig2 = px.choropleth(df2, locations='state', locationmode='USA-states',
                         color='is_fraud', scope='usa',
                         title='Fraud Frequency by State')
    fig2.show()

if 'category' in df2.columns:
    fig3 = px.bar(df2.groupby('category')['is_fraud'].mean().reset_index(),
                  x='category', y='is_fraud', title='Average Fraud Rate by Category')
    fig3.show()
```

PDF Report Generation

- Creates a professionally formatted PDF summarizing project results
 - *reportlab* allowed for structured document creation
- Includes contents generated from these models
 - Title and timestamp
 - Dataset overview
 - Model performance
 - ROC AUC scores for all tested models
 - Key insights:
 - RandomForest was the top performer
 - SHAP explained key fraud indicators
 - Visual dashboards highlight fraud trends by amount, category, and location

```
# === Cell 18: Generate Summary PDF Report ===
from reportlab.lib.pagesizes import letter
from reportlab.platypus import SimpleDocTemplate, Paragraph, Spacer
from reportlab.lib.styles import getSampleStyleSheet
import datetime

summary_path = "CreditCard_Fraud_Report.pdf"
doc = SimpleDocTemplate(summary_path, pagesize=letter)
styles = getSampleStyleSheet()
story = []

story.append(Paragraph("<b>Credit Card Fraud Detection Report</b>", styles['Title']))
story.append(Spacer(1, 12))
story.append(Paragraph(f"Generated on: {datetime.datetime.now()}", styles['Normal']))
story.append(Spacer(1, 12))

# Add dataset info
story.append(Paragraph(f"Rows: {df2.shape[0]}, Columns: {df2.shape[1]}", styles['Normal']))
story.append(Spacer(1, 12))

# Add model results
if 'results' in locals():
    story.append(Paragraph("<b>Model Comparison (ROC AUC)</b>", styles['Heading2']))
    for _, r in pd.DataFrame(results).iterrows():
        story.append(Paragraph(f"{'r['Model']}: {round(r['ROC_AUC'], 3)}", styles['Normal']))
    story.append(Spacer(1, 6))

# Add key insights
story.append(Paragraph("<b>Key Insights:</b>", styles['Heading2']))
story.append(Paragraph("• RandomForest performed best among tested models.", styles['Normal']))
story.append(Paragraph("• SHAP analysis shows top drivers of fraud risk.", styles['Normal']))
story.append(Paragraph("• Interactive charts visualize fraud by amount, category, and state.", styles['Normal']))

doc.build(story)
print(f"✅ Report generated: {summary_path}")
```

AI Aspects to be Included

- LLM integration to explain the features causing fraud detections using prompts.
- Explain the insights produced by models using LLM
- Agents that simulate fraud data to decrease the False Positive Rates.

Next Steps

- Creation and utilization of API keys for FRED and OpenAI
 - Integrate economic elements into model analysis
- Additional generative AI aspects?
 - Expand on insights that can be produced from fraudulent detection model
- Being able to segment data into smaller sets
 - Potentially limited generation of certain arrays due to the demand on local computer RAM

References

- [A systematic review of AI-enhanced techniques in credit card fraud detection | Journal of Big Data](#)
- [Generative AI in Card Fraud Detection: Benefits, Use Cases, and Industry Impact](#)
- [FraudX AI: An Interpretable Machine Learning Framework for Credit Card Fraud Detection on Imbalanced Datasets](#)
- [AI Analytics for Credit Card Security and Fraud Prevention](#)